Double Descent in Linear Models

Andrea Vaselli
Statistical Methods for Machine Learning, A.Y. 2025/26
Code: `https://github.com/andreavase27/double-descent-linear-models`

Abstract

Following Belkin et al., we study the *double descent* curve in linear regression, fixing the number of training points ($n = 100$) and varying the number of features d. We compare minimum-norm least squares and ridge regression on synthetic data with noisy linear labels and on the California Housing dataset. In every setting the test error peaks exactly at the interpolation threshold $d = n$, together with the norm of the coefficients, and decreases again beyond it; on synthetic data this matches a closed-form expression for the expected error. Whether overparameterized models beat the classical ones depends on how the signal is spread across the features, and ridge regression with a cross-validated penalty removes the peak. As extensions, we show that gradient descent started at zero converges to the minimum-norm solution and that the peak grows with the label noise.

1 Introduction

The classical view of generalization is built on the bias–variance trade-off. A model that is too simple cannot capture the structure of the data (high bias), while a model that is too complex starts fitting the noise (high variance), so the test error follows a U-shaped curve. A direct consequence is that a model with zero training error, one that interpolates the data, should generalize poorly.

Modern machine learning does not seem to follow this rule. Large models, such as neural networks with far more parameters than training points, are routinely trained to zero training error and still perform well on new data. Belkin et al. [1] proposed a way to reconcile the two pictures: the *double descent* curve. The U-shaped curve is only its first part. At the interpolation threshold, where the model has just enough parameters to fit the training data exactly, the test error reaches a peak; beyond it, adding parameters makes the test error go down again, sometimes below the best value of the classical regime.

The goal of this project is to reproduce this phenomenon in the simplest possible setting, linear regression. We fix the number of training points n, vary the number of features d, and compare minimum-norm least squares with ridge regression, on a synthetic dataset with noisy linear labels and on the California Housing dataset. As extensions, we compare gradient descent with the closed-form solution and study the effect of the label noise.

We find the double descent curve in every setting, always with the peak at $d = n$. Whether the second descent beats the classical regime, however, depends on how the signal is spread across the features, and ridge regression with a well-chosen penalty removes the peak altogether.

2 Background

2.1 Least squares and the interpolation threshold

We have n training points, collected in a matrix $\mathbf{X} \in \mathbb{R}^{n \times d}$ (one row per point, one column per feature) and a vector of labels $\mathbf{y} \in \mathbb{R}^n$. Least squares looks for the coefficients $\mathbf{w}$ that minimize

the training error $\|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2$. Setting the gradient to zero gives the normal equations

$$\mathbf{X}^\top \mathbf{X}\, \mathbf{w} = \mathbf{X}^\top \mathbf{y}. \tag{1}$$

This is a system with one equation per training point and one unknown per feature, so what happens depends on how d compares with n. When $d < n$ there are more equations than unknowns: in general no $\mathbf{w}$ fits all the points, and least squares picks the one that comes closest, $\hat{\mathbf{w}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$. When $d = n$ there is exactly one solution, and it fits every training point: the training error is zero. This is the *interpolation threshold*. When $d > n$ there are infinitely many solutions, all with zero training error, and we have to choose one.

We choose the one with the smallest Euclidean norm. Any $\mathbf{w}$ can be split into a part that is a combination of the rows of $\mathbf{X}$ and a part that $\mathbf{X}$ sends to zero. The second part does not change any prediction on the training set, but it makes $\mathbf{w}$ longer, so the shortest solution has it equal to zero and is of the form $\mathbf{w} = \mathbf{X}^\top \mathbf{a}$. Imposing $\mathbf{X}\mathbf{w} = \mathbf{y}$ gives

$$\hat{\mathbf{w}} = \mathbf{X}^\top (\mathbf{X}\mathbf{X}^\top)^{-1} \mathbf{y} = \mathbf{X}^+ \mathbf{y}, \tag{2}$$

where $\mathbf{X}^+$ is the Moore–Penrose pseudoinverse. Both cases are therefore the same estimator, $\hat{\mathbf{w}} = \mathbf{X}^+ \mathbf{y}$.

2.2 Why the test error peaks at $d = n$

Close to $d = n$ the data are just barely enough to determine all the coefficients, and the matrix $\mathbf{X}^\top \mathbf{X}$ (or $\mathbf{X}\mathbf{X}^\top$) is close to singular: there is at least one direction in which the training points vary very little. To pass exactly through every point, noise included, the solution needs very large coefficients along that direction. The norm of $\hat{\mathbf{w}}$ explodes, and so does the variance of the predictions on new data. Beyond the threshold every extra feature gives new directions in which the noise can be absorbed: the minimum-norm interpolant becomes shorter and the test error goes down again. This is the second descent.

2.3 Ridge regression

Ridge regression adds a penalty on the size of the coefficients:

$$\hat{\mathbf{w}}_\lambda = \arg\min_{\mathbf{w}} \; \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2 + \lambda \|\mathbf{w}\|^2 = (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^\top \mathbf{y} = \mathbf{X}^\top (\mathbf{X}\mathbf{X}^\top + \lambda \mathbf{I})^{-1} \mathbf{y}. \tag{3}$$

For $\lambda > 0$ the matrix to invert is always invertible, whatever d is. The two expressions give the same result: the first (primal form) solves a $d \times d$ system, the second (dual form) an $n \times n$ one. In the directions where the data vary little, the term λ keeps the coefficients from growing, which is exactly what lowers the peak. As $\lambda \to 0$, ridge recovers the minimum-norm solution.

2.4 Expected test error with Gaussian features

For our synthetic model, with Gaussian features and labels $y = \mathbf{x}^\top \boldsymbol{\beta} + \sigma\varepsilon$, the expected test error of minimum-norm least squares that uses the first d features has a closed form [2]. Let s_d be the share of the signal carried by the first d features, and $\tau^2 = \sigma^2 + (1 - s_d)$ the *effective noise*: the label noise plus the signal the model cannot see. Then

$$\mathbb{E}[\text{test error}] = \begin{cases} \tau^2 \left(1 + \dfrac{d}{n - d - 1}\right) & d < n - 1, \\ s_d \left(1 - \dfrac{n}{d}\right) + \tau^2 \left(1 + \dfrac{n}{d - n - 1}\right) & d > n + 1, \end{cases} \tag{4}$$

and it is infinite for $n - 1 \le d \le n + 1$. Below the threshold the effective noise is amplified by a factor that blows up as d approaches n. Above it there are two terms: the part of the visible signal that n points are not enough to recover, $s_d(1 - n/d)$, and the effective noise, amplified by a factor that goes back towards 1 as d grows.

3 Datasets

3.1 Synthetic data

In order to obtain the synthetic dataset we draw inputs $\mathbf{x} \in \mathbb{R}^D$ with $D = 500$ independent standard Gaussian entries, and generate the labels as

$$y = \mathbf{x}^\top \boldsymbol{\beta} + \sigma\varepsilon, \qquad \varepsilon \sim \mathcal{N}(0,1), \quad \sigma = 0.5. \tag{5}$$

The vector $\boldsymbol{\beta}$ is drawn at random and then rescaled so that $\|\boldsymbol{\beta}\| = 1$. Since the features are independent and have unit variance, the signal $\mathbf{x}^\top\boldsymbol{\beta}$ has variance $\|\boldsymbol{\beta}\|^2 = 1$: the signal-to-noise ratio is $1/\sigma^2 = 4$ and is controlled by σ alone. A model of size d only sees the first d features. The signal carried by the other $D - d$ features cannot be explained by the model, so it acts as extra noise.

In every repetition we draw $n = 100$ training points and 2000 test points independently from the same distribution. The two sets never overlap, so there is no leakage, and with 2000 test points the estimate of the test error is precise. We take $D = 5n$ so that there is enough room after the interpolation threshold $d = n$ to see the second descent. Two reference values are useful when reading the plots: no predictor can go below the noise variance $\sigma^2 = 0.25$, and the trivial predictor that always outputs 0 has error $1 + \sigma^2 = 1.25$.

We also want to see how the way the signal is spread across the features changes the curve. To do so, before normalizing $\boldsymbol{\beta}$ we multiply its j-th entry by j^{-a}. In *scenario A* ($a = 0$) all features are equally important on average, so each new feature adds only a small piece of the signal. In *scenario B* ($a = 1$) most of the signal is in the first features, as often happens with real data.

3.2 Real data: California Housing

As the real-world dataset we use California Housing [3], loaded with scikit-learn [4]. It is made of 20,640 observations, which are the districts of California, and 8 numeric features, which are the following: median income, house age, average number of rooms and bedrooms, population, average occupancy, latitude and longitude. The dependent variable is the median house value of the district, expressed in units of 100,000 dollars. We picked this dataset because it is a regression problem on real data and it is large enough to draw a new training set and test set in every repetition. Its limitation is that the target is capped at 5 (500,000 dollars), so the most expensive districts all have the same value.

In every repetition we shuffle the districts and take the first 100 as training set and the next 2000 as test set, so the two sets never overlap. Inputs and target are standardized using the mean and standard deviation of the training set only, and the same transformation is then applied to the test set. Indeed, computing these statistics on the whole dataset would leak information from the test set into training. After standardization, predicting the training mean gives a test error of about 1, which is the reference value to beat.

With only 8 features the model could never reach the interpolation threshold $d = n = 100$. As in [1], we therefore map each input to random Fourier features [5]: we draw 1000 random vectors $\mathbf{v}_k \sim \mathcal{N}(\mathbf{0}, I/h^2)$ and compute $\cos(\mathbf{v}_k^\top \mathbf{x})$ and $\sin(\mathbf{v}_k^\top \mathbf{x})$, which gives up to $d = 2000$ features. Here only the inputs change; the model is still linear in the coefficients. The columns are ordered in pairs, so that the first d columns always contain complete cosine–sine pairs and d is always even. The same random vectors are used for both the training and the test set. To set the oscillation speed of the features, we fix a bandwidth h. It is chosen with the median heuristic: h is the median distance between pairs of training points, so that it never depends on the test set. Since this choice turns out to matter, we also report results for h divided by 2 and by 4.

4 Experimental protocol

Model sizes. The number of training points is fixed at $n = 100$ and the number of features d is varied. Far from the threshold the grid of values of d is sparse, while around the threshold, where the error changes quickly, it is dense. Specifically, we use steps of 5 up to $d = 75$, steps of 2 between 80 and 120, then steps of 10 up to 200 and of 25 up to 500 on synthetic data. On California Housing the grid is the same around the threshold but goes up to $d = 2000$, and only even values are used. The gradient descent experiment, which is slower, uses a coarser grid.

Repetitions and summary statistics. The training set is random, and so is the test error of a model trained on it. Each experiment is therefore repeated 50 times, using seeds 0 to 49: on synthetic data every repetition generates new training and test sets, while on California Housing it draws a new random split of the same pool of districts, so the repetitions are nearly but not exactly independent. For every d we report the median over the repetitions and the interquartile range (from the 25th to the 75th percentile). We use the median rather than the mean because, by Equation (4), the expected test error is infinite around $d = n$: close to the threshold a few unlucky repetitions produce huge errors, and the mean would depend almost only on them.

Metrics. For each model we record the training and test mean squared errors and the norm $\|\hat{\mathbf{w}}\|$ of the estimated coefficients, which explains the behaviour of the test error near the threshold.

Models and choice of λ. We compare minimum-norm least squares with ridge regression for three fixed penalties, $\lambda \in \{1, 10, 100\}$, which show the effect of increasing regularization, and with ridge regression where λ is chosen by 5-fold cross-validation among 13 values evenly spaced on a log scale between 10^{-3} and 10^{3}. To choose λ we only look at the training set. We split the 100 training points into 5 folds of 20 points. For each candidate value of λ we train the model 5 times, each time leaving out a different fold and measuring the error on it, and we average the 5 errors. We then keep the λ with the lowest average error and retrain the model on all 100 points. The test set is used only at the end, to measure the error of the final model.

Implementation. All algorithms are implemented from scratch in Python; NumPy is used only for basic linear algebra. Least squares and ridge solve a linear system with the function `numpy.linalg.solve` rather than computing a matrix inverse, because this is more accurate when the matrix is close to singular, as happens near $d = n$. When $d \leq n$ they both use the primal form ($d \times d$ system), when $d > n$ they both use the dual form ($n \times n$ system). We checked the implementation in two ways: the least squares coefficients match those obtained with the pseudoinverse up to differences of order 10^{-15} (about 10^{-11} at $d = n$, where the system is close to singular), and the simulated test error on synthetic data matches the theoretical curve of Equation (4). Gradient descent always starts from $\mathbf{w} = \mathbf{0}$, which is what makes it converge to the minimum-norm solution. It uses a fixed step size $1/L$, where L is the largest eigenvalue of $\mathbf{X}^\top \mathbf{X}/n$: this is the standard choice that guarantees the training error decreases at every step. Since gradient descent is much slower than the closed form, in this experiment we use 20 repetitions instead of 50.

Visualizations. In all figures a vertical line marks the threshold $d = n$, curves are medians and shaded bands are interquartile ranges. Errors and norms are on a logarithmic scale, which becomes linear below 0.1 for the training error so that zero can be shown. Dotted lines are reference values (noise floor, trivial predictor) and dashed lines the theoretical curve.

Reproducibility. The random draws use a NumPy generator with a fixed seed, and the repository contains the library versions (`requirements.txt`) and one script per experiment, so every figure in this report can be regenerated with a single command.

5 Results

5.1 Synthetic data: least squares

Figure 1 shows the results of minimum-norm least squares on the two synthetic scenarios. In both of them the curve has the same shape. The training error goes down as d grows and becomes exactly zero at $d = 100$: from that point on the model interpolates the training data. At the same point the test error explodes, reaching a median of about 278 in scenario A and 49 in scenario B, and then decreases again as d keeps growing. The interpolation threshold is therefore exactly at $d = n$, as expected, and the norm of $\hat{\mathbf{w}}$ peaks in the same place.

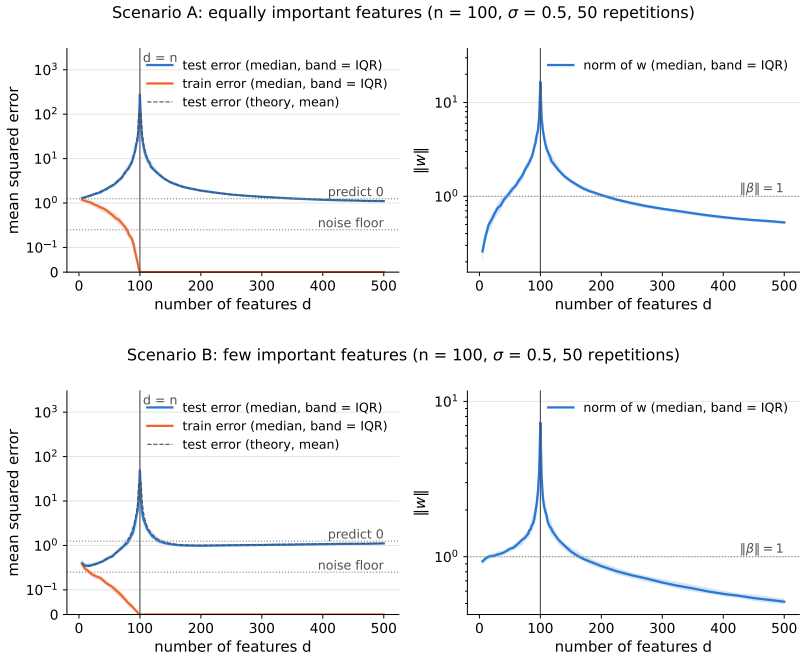


Figure 1: Minimum-norm least squares on synthetic data (top: scenario A, bottom: scenario B). Left: median train and test error over 50 repetitions, with interquartile band and the theoretical expected error (dashed). Right: norm of the estimated coefficients. The vertical line marks the interpolation threshold $d = n = 100$.

The theoretical curves of Equation (4) follow the simulations closely in both scenarios, which confirms the implementation. The two scenarios differ in what happens away from the threshold. In scenario A there is no classical U-shape: the test error grows from the start, because each new feature adds very little signal and some variance. After the threshold it keeps decreasing, and the best model is the largest one, $d = 500$, with a median error of 1.10. This is the main claim of the paper: a model with five times more parameters than data points does better than any model below the threshold. The gain, however, is small: the trivial predictor that always outputs 0 has error 1.25. In scenario B the U-shape appears: the error first goes down, reaching its minimum of 0.35 at $d = 15$, and then goes up. The second descent is still there, but it never comes back below this minimum.

5.2 Synthetic data: comparison with ridge

Figure 2 compares least squares with ridge regression. The larger λ, the lower the peak: in scenario A the height of the peak (median test error) goes from 278 with least squares to 6.2 with $\lambda = 1$, 2.4 with $\lambda = 10$ and 1.3 with $\lambda = 100$. The norm of the coefficients, shown on the right, no longer explodes. When λ is chosen by cross-validation the peak disappears completely and the test error stays almost flat as d grows. In our experiments, then, double descent is a property of the unregularized estimator: with a well-chosen penalty there is no peak at all.

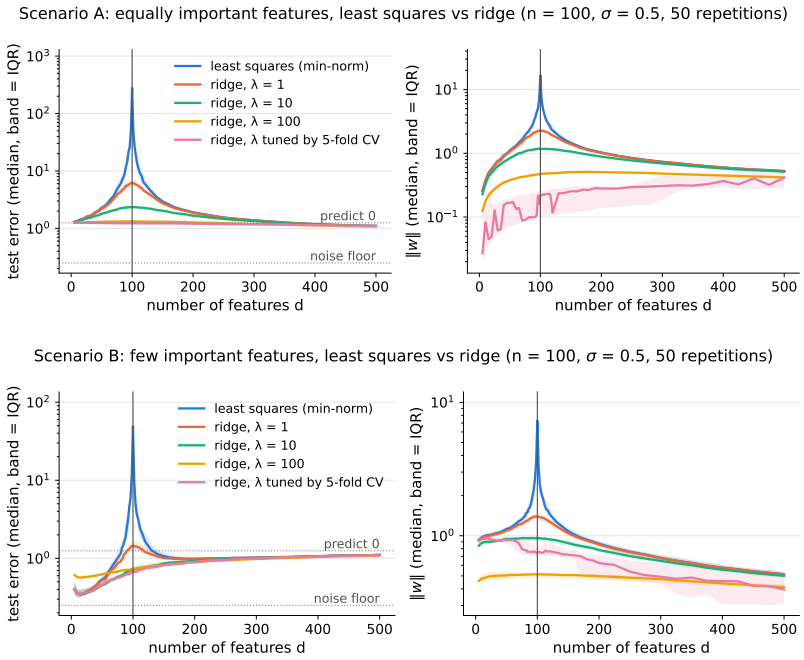


Figure 2: Least squares versus ridge regression on synthetic data (top: scenario A, bottom: scenario B). Left: median test error; right: norm of the coefficients. Ridge with λ tuned by 5-fold cross-validation on the training set removes the peak.

Far from the threshold all the methods give almost the same error (about 1.1 at $d = 500$ in scenario A): there the minimum-norm solution already behaves like a regularized one. A penalty that is too strong also has a cost. In scenario B, where a few features carry most of the signal, $\lambda = 100$ shrinks the important coefficients as well, and the best error rises from 0.35 to 0.57. With λ chosen by cross-validation, instead, the best error is 0.34, the lowest of all. In scenario A the cross-validated model stays close to the trivial predictor: with 100 points and 500 weak features there is little to learn, and cross-validation detects it.

5.3 Real data: California Housing

On California Housing the curve has the same shape as on synthetic data (Figure 3). The training error reaches zero at $d = 100$, the test error has a sharp peak at the same point, with a median of about 4800, and the norm of the coefficients peaks there too. After the threshold the test error decreases again. Double descent therefore appears on real data as well, with the threshold exactly at $d = n$.

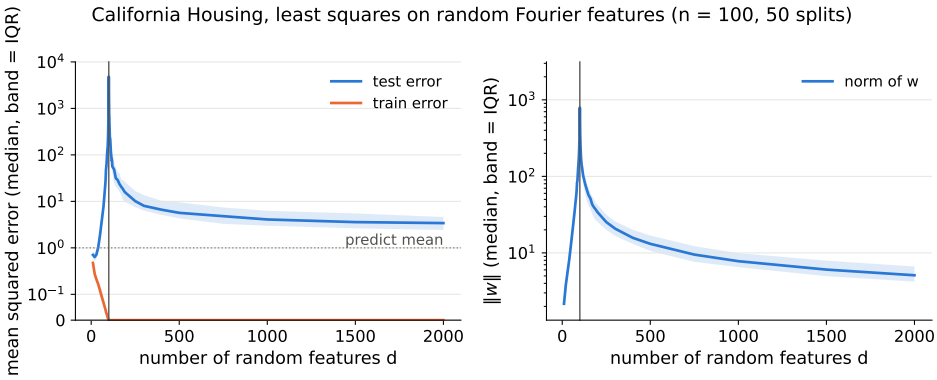


Figure 3: California Housing, minimum-norm least squares on random Fourier features (bandwidth from the median heuristic), 50 random train/test splits. Left: train and test error; right: norm of the coefficients.

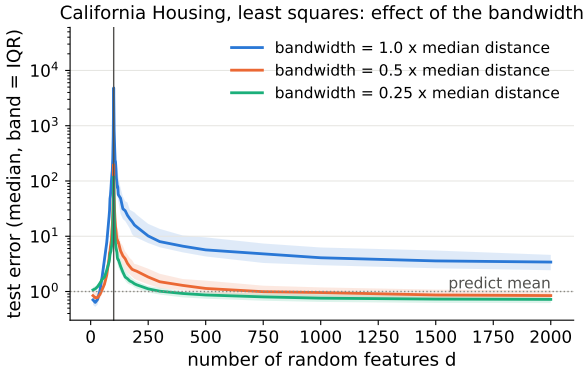


Figure 4: California Housing, least squares with three bandwidths of the random features. Narrower features make the overparameterized regime competitive.

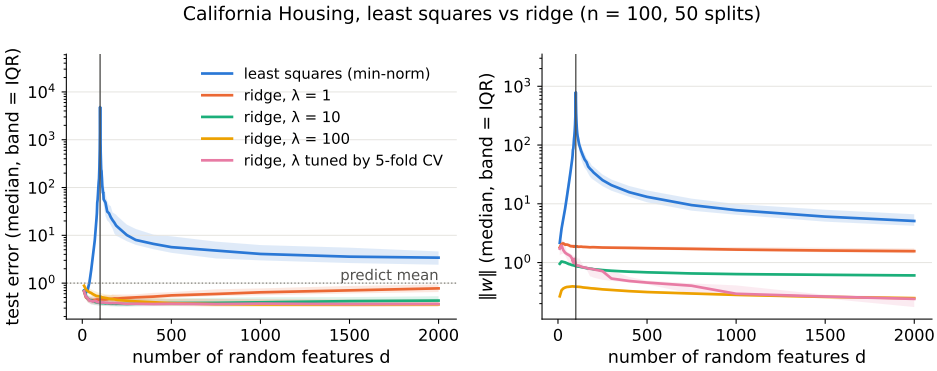


Figure 5: California Housing, least squares versus ridge regression on the same random features.

With the bandwidth chosen by the median heuristic, however, the second descent is not enough. The best model is a small one, $d = 20$, with a median error of 0.63, while at $d = 2000$ the error is still around 3, worse than predicting the mean. This is the same behaviour as scenario B. Figure 4 shows that the bandwidth changes this conclusion. With narrower features (h divided by 4) the peak is much lower and the error keeps decreasing up to $d = 2000$, where it reaches 0.72, below the trivial predictor and better than every model below the threshold with

the same bandwidth. Whether overparameterization pays off on this dataset thus depends on how the features are built.

Ridge regression behaves in a similar way to synthetic data (Figure 5). With $\lambda \geq 1$ the peak disappears, and with λ chosen by cross-validation the error keeps improving as features are added, down to 0.36 at $d = 1500$; ridge with $\lambda = 100$ does about as well (0.357 at $d = 1000$). The best results on this dataset are therefore reached by strongly overparameterized models with the right amount of regularization.

5.4 Extension: gradient descent versus closed form

Figure 6 compares gradient descent started at $\mathbf{w} = \mathbf{0}$ with the closed-form solution. Away from the threshold the two coincide: after 1000 iterations the relative distance between them is of order 10^{-15}, both for $d < n$ and for $d > n$. So for $d > n$ gradient descent finds exactly the minimum-norm interpolant, even though nothing in the algorithm asks for it. This is the implicit regularization of gradient descent: since every update is a combination of the rows of $\mathbf{X}$, the iterates never leave the subspace that contains the minimum-norm solution.

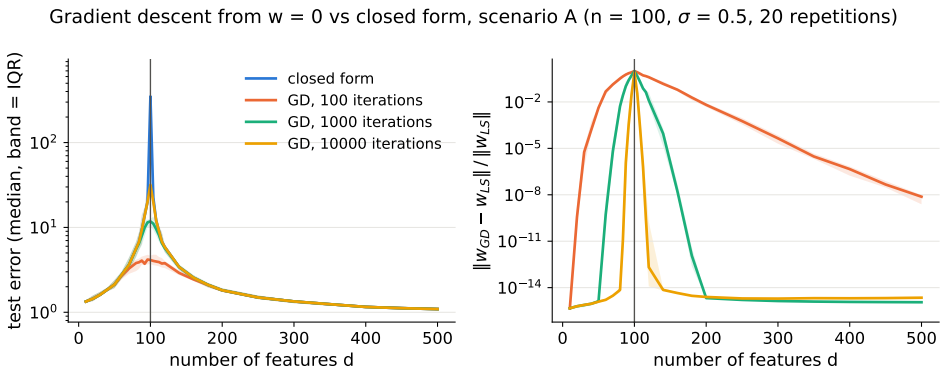


Figure 6: Gradient descent started at $\mathbf{w} = \mathbf{0}$ versus the closed-form minimum-norm solution (scenario A, 20 repetitions). Left: median test error for three iteration budgets. Right: relative distance between the two solutions.

Close to the threshold gradient descent does not converge in practice: at $d = 100$ the relative distance is still 0.83 after 10,000 iterations. Convergence is slow there for the same reason the peak exists: the smallest nonzero eigenvalue of $\mathbf{X}^\top\mathbf{X}$ is close to zero, and along that direction each step moves very little. As a consequence the peak of gradient descent is lower, and it grows with the number of iterations: the median error at the threshold is about 4 after 100 iterations, about 12 after 1000 and about 30 after 10,000, against roughly 350 for the closed form. Stopping early acts like ridge regression: it prevents the coefficients from growing along the directions where they would explode.

5.5 Extension: effect of the label noise

Figure 7 repeats the experiment on synthetic data with four noise levels. More noise gives a higher peak in both scenarios: in scenario B the median error at $d = n$ goes from 1.3 without noise to 15, 49 and 196 for $\sigma = 0.25$, 0.5 and 1. Again, the theoretical curves match the simulations.

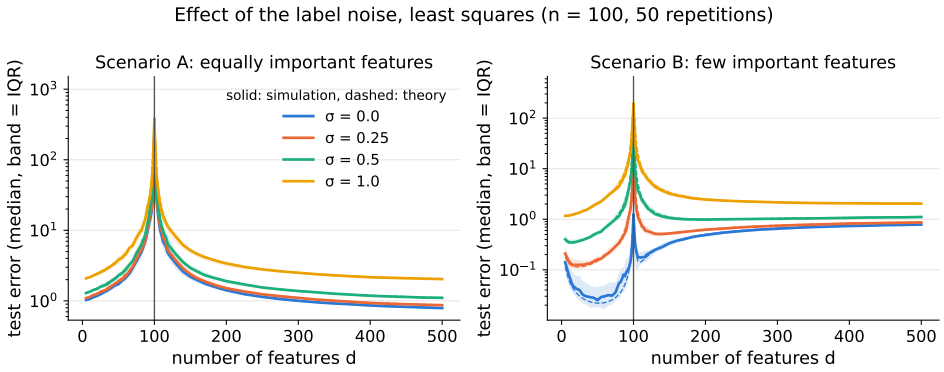


Figure 7: Test error of minimum-norm least squares for four noise levels, simulation (solid) and theory (dashed).

The two scenarios behave very differently when there is no noise. In scenario B the peak almost disappears, because the first 100 features contain nearly all the signal and there is nothing left to fit by mistake. In scenario A, instead, the peak is still large, about 179, even with $\sigma = 0$. The reason is that a model with $d = 100$ sees only a fifth of the signal, and the rest acts as noise. Double descent therefore does not require noise in the labels: it is enough that the model cannot explain everything.

Noise also moves the best model. In scenario B the best d goes from 50 without noise to 15 for $\sigma = 0.25$ and $\sigma = 0.5$, and to 5 for $\sigma = 1$: the noisier the data, the smaller the model that works best. Below the threshold the classical bias-variance trade-off still holds.

6 Discussion

Where the classical picture fails. All our experiments show the same thing: beyond the interpolation threshold, adding parameters does not make the model overfit more. The test error goes down again, and in some settings (scenario A, California Housing with narrow features) the best model of all has many more parameters than data points. The classical view fails because it measures complexity by the number of parameters. What actually drives the test error is the norm of the selected solution: it peaks at $d = n$ exactly where the test error does, and it decreases afterwards. The minimum-norm choice, whether made explicitly or reached implicitly by gradient descent started at zero, acts as a regularizer that the parameter count does not see.

Strengths. The threshold is always at $d = n$, on every dataset, noise level and bandwidth, the simulations match the theory, and both extensions are consistent with the main picture.

Limitations. The results also have clear limits. In scenario A the overparameterized model wins, but only barely: its error is 1.10, against 1.25 for the trivial predictor. In scenario B and on California Housing with the median heuristic, the second descent never goes below the classical minimum, so overparameterization is not always worth it. Moreover, ridge regression with a cross-validated penalty removes the peak and is among the best models in every setting: double descent is mostly a property of unregularized estimators, and with a well-tuned regularizer the classical concern about the peak disappears. Finally, our setting is narrow: a single training size $n = 100$, Gaussian features in the synthetic experiments, a target that is capped in California Housing, and on real data a dependence on the bandwidth of the random features, which we fixed with a heuristic rather than tuned.

Synthetic versus real data. The two datasets agree on the position of the peak, the behaviour of the norm and the effect of ridge. They also agree on when overparameterization helps: California Housing with the median heuristic behaves like scenario B, and with narrower features like scenario A.

7 Conclusion

We reproduced the double descent curve for linear regression on synthetic and real data, with the peak always at $d = n$. Overparameterized models can beat the classical ones, but only when the signal is spread across many features or the features are built appropriately, and a well-chosen ridge penalty removes the peak. Double descent does not overturn the bias–variance trade-off, but it shows that the number of parameters is not the right measure of complexity: the norm of the solution is.

References

[1] Mikhail Belkin, Daniel Hsu, Siyuan Ma, and Soumik Mandal. Reconciling modern machine-learning practice and the classical bias–variance trade-off. *Proceedings of the National Academy of Sciences*, 116(32):15849–15854, 2019.

[2] Mikhail Belkin, Daniel Hsu, and Ji Xu. Two models of double descent for weak features. *SIAM Journal on Mathematics of Data Science*, 2(4):1167–1180, 2020.

[3] R. Kelley Pace and Ronald Barry. Sparse spatial autoregressions. *Statistics & Probability Letters*, 33(3):291–297, 1997.

[4] Fabian Pedregosa et al. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.

[5] Ali Rahimi and Benjamin Recht. Random features for large-scale kernel machines. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 20, 2007.

Declaration

I declare that this material, which I now submit for assessment, is entirely my own work and has not been taken from the work of others, save and to the extent that such work has been cited and acknowledged within the text of my work. I understand that plagiarism, collusion, and copying are grave and serious offences in the university and accept the penalties that would be imposed should I engage in plagiarism, collusion or copying. This assignment, or any part of it, has not been previously submitted by me or any other person for assessment on this or any other course of study.